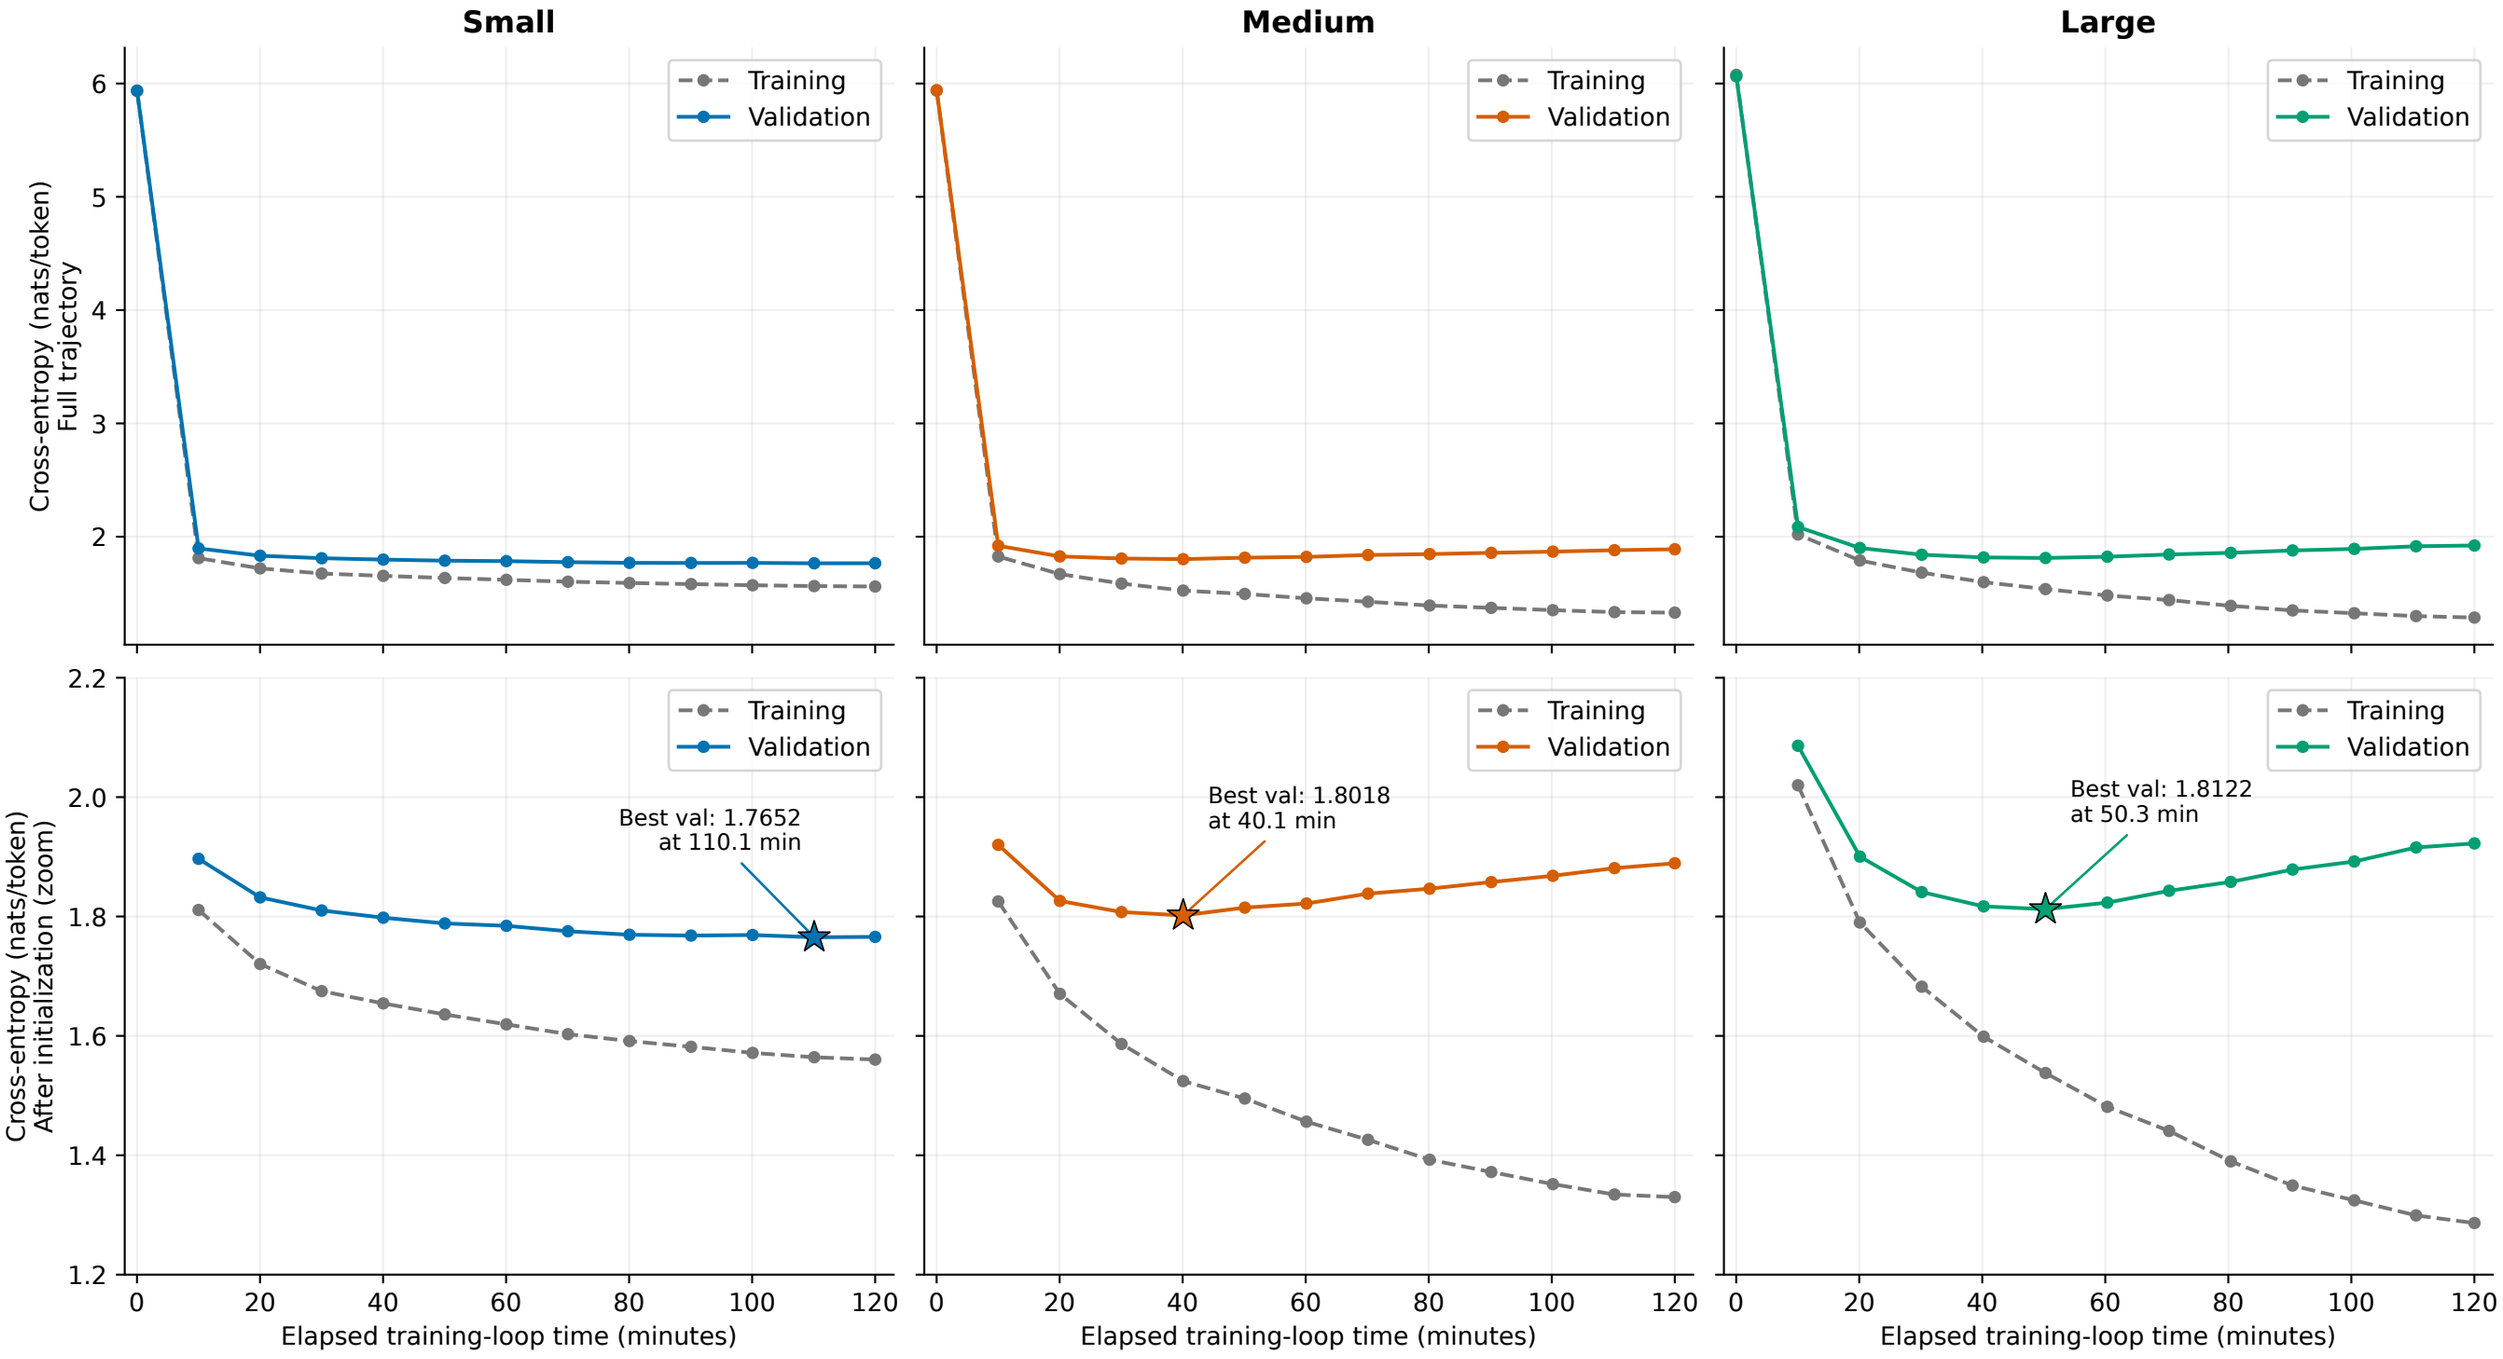
MAESTRO: training and validation loss at saved evaluation checkpoints
Top: all 13 evaluations. Bottom: 12 evaluations after the untrained baseline. Lines connect recorded points.
Small
Medium
Large
Training
Validation
Cross-entropy (nats/token)
Full trajectory
Cross-entropy (nats/token)
After initialization (zoom)
Best val: 1.7652
at 110.1 min
Best val: 1.8018
at 40.1 min
Best val: 1.8122
at 50.3 min
Elapsed training-loop time (minutes)